

In [3]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [4]:

```
df=pd.read_csv(r"C:\Users\sruth\Downloads\bottle.csv.zip")  
df
```

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\1331323283.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.

```
df=pd.read_csv(r"C:\Users\sruth\Downloads\bottle.csv.zip")
```

Out[4]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600

In [5]:

```
df=df[['Salnty','T_degC']]
2      1      3      054.0      19-4903CR-HY-060-0930-05400560-0010A-7
056.0      10      10.460      33.4370      NaN      25.65400
```

In [6]:

```
df.columns=['Sal','Temp']
```

In [7]:

```
df.head(10)
```


Out[7]:

	4	1		5	054.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10.450	33.4210	NaN	25.64300
	Sal	Temp			056.0						
0	33.440	10.50									
1	33.440	10.46
2	33.437	10.46									
3	33.420	10.45				20-1611SR-1611SR-2239-09340264-0000A-7					
864858	34404	34404	864859	093.4	026.4	MX-310-2239-09340264-0002A-3	0	18.744	33.4083	5.805	23.87055
4	33.421	10.45									
5	33.431	10.45									
6	33.440	10.45									
864859	34404	34404	864860	093.4	026.4	20-1611SR-1611SR-2239-09340264-0002A-3	2	18.744	33.4083	5.805	23.87072
8	33.420	10.06									
9	33.494	9.86									
864860	34404	34404	864861	093.4	026.4	20-1611SR-1611SR-2239-09340264-0005A-3	5	18.692	33.4150	5.796	23.88911
864861	34404	34404	864862	093.4	026.4	20-1611SR-1611SR-2239-09340264-0010A-3	10	18.161	33.4062	5.816	24.01426

```
In [8]: Cst_Cnt Btl_Cnt Sta_ID Depth_ID Depthm T_degC Salnty O2ml_L STheta
```

```
sns.lmplot(x="Sal",y="Temp",data=df,order=2,ci=None)
```

```
Out[8]: 20-  
1611SR-  
MX-310-  
15 17.533 33.3880 5.774 24.15297  
<seaborn.axisgrid.FacetGrid at 0x224521654550>  
09340264-
```



```
In [9]:
```

```
df.describe()
```

```
Out[9]:
```

	Sal	Temp
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Sal      817509 non-null  float64
 1   Temp     853900 non-null  float64
dtypes: float64(2)
memory usage: 13.2 MB
```

In [11]:

```
df.fillna(method='ffill',inplace=True)
```

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\3337295870.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(method='ffill',inplace=True)

In [12]:

```
x=np.array(df['Sal']).reshape(-1,1)
y=np.array(df['Temp']).reshape(-1,1)
```

In [13]:

```
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


0.20304614002564747

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\693062840.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.dropna(inplace=True)


In [14]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_pred,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [15]:

```
df500=df[:][:500]
sns.lmplot(x="Sal",y="Temp",data=df500,order=1,ci=None)
x=np.array(df500['Sal']).reshape(-1,1)
y=np.array(df500['Temp']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```




In [16]:

```
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
```

Regression: 0.8627231173560658

In [17]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [18]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: 0.8627231173560658

In [19]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```


In [20]:

```
df=pd.read_csv(r"C:\Users\sruth\Downloads\fiat500_VehicleSelection_Dataset.csv")
df
```

Out[20]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611111
1	2	pop	51	1186	32500	1	45.666359	12.244444
2	3	sport	74	4658	142228	1	45.503300	11.411111
3	4	lounge	51	2739	160000	1	40.633171	17.633333
4	5	pop	73	3074	106880	1	41.903221	12.494444
...
1533	1534	sport	51	3712	115280	1	45.069679	7.700000
1534	1535	lounge	74	3835	112000	1	45.845692	8.666667
1535	1536	pop	51	2223	60457	1	45.481541	9.411111
1536	1537	lounge	51	2557	80750	1	45.000702	7.688889
1537	1538	pop	51	1766	54276	1	40.323410	17.566667

1538 rows × 9 columns

In [21]:

```
df=df[['age_in_days','km']]
df.columns=['age','km']
```

In [22]:

```
df.head(10)
```

Out[22]:


	age	km
0	882	25000
1	1186	32500
2	4658	142228
3	2739	160000
4	3074	106880
5	3623	70225
6	731	11600
7	1521	49076
8	4049	76000
9	3653	89000

In [23]:

```
sns.lmplot(x="age",y="km",data=df,order=2,ci=None)
```

Out[23]:

<seaborn.axisgrid.FacetGrid at 0x2465d2b2080>



In [24]:

```
df.describe()
```

Out[24]:

	age	km
count	1538.000000	1538.000000
mean	1650.980494	53396.011704
std	1289.522278	40046.830723
min	366.000000	1232.000000
25%	670.000000	20006.250000
50%	1035.000000	39031.000000
75%	2616.000000	79667.750000
max	4658.000000	235000.000000

In [25]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   age      1538 non-null    int64  
 1   km       1538 non-null    int64  
dtypes: int64(2)
memory usage: 24.2 KB
```

In [26]:

```
df.fillna(method='ffill',inplace=True)
```

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\3337295870.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(method='ffill',inplace=True)

In [27]:

```
x=np.array(df['age']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
```

In [28]:

```
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


```
0.7099400299873513
```

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\693062840.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.dropna(inplace=True)


In [29]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_pred,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [30]:

```
df500=df[:][:500]
sns.lmplot(x="age",y="km",data=df500,order=1,ci=None)
x=np.array(df500['age']).reshape(-1,1)
y=np.array(df500['km']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```




In [31]:

```
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
```

Regression: 0.647134099766666

In [32]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [33]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```


In [34]:

```
df=pd.read_csv(r"C:\Users\sruth\Downloads\data.csv")
df
```

Out[34]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	2014-05-02 00:00:00	3.130000e+05	3.0	1.50	1340	7912	1.5	0
1	2014-05-02 00:00:00	2.384000e+06	5.0	2.50	3650	9050	2.0	0
2	2014-05-02 00:00:00	3.420000e+05	3.0	2.00	1930	11947	1.0	0
3	2014-05-02 00:00:00	4.200000e+05	3.0	2.25	2000	8030	1.0	0
4	2014-05-02 00:00:00	5.500000e+05	4.0	2.50	1940	10500	1.0	0
...
4595	2014-07-09 00:00:00	3.081667e+05	3.0	1.75	1510	6360	1.0	0
4596	2014-07-09 00:00:00	5.343333e+05	3.0	2.50	1460	7573	2.0	0
4597	2014-07-09 00:00:00	4.169042e+05	3.0	2.50	3010	7014	2.0	0
4598	2014-07-10 00:00:00	2.034000e+05	4.0	2.00	2090	6630	1.0	0
4599	2014-07-10 00:00:00	2.206000e+05	3.0	2.50	1490	8102	2.0	0

4600 rows × 18 columns



In [35]:

```
df=df[['sqft_living','sqft_lot']]  
df.head(10)
```

Out[35]:


	sqft_living	sqft_lot
0	1340	7912
1	3650	9050
2	1930	11947
3	2000	8030
4	1940	10500
5	880	6380
6	1350	2560
7	2710	35868
8	2430	88426
9	1520	6200

In [36]:

```
sns.lmplot(x='sqft_living',y='sqft_lot',data=df,order=2,ci=None)
```

Out[36]:

<seaborn.axisgrid.FacetGrid at 0x2465cb832b0>



In [37]:

```
df.describe()
```

Out[37]:

	sqft_living	sqft_lot
count	4600.000000	4.600000e+03
mean	2139.346957	1.485252e+04
std	963.206916	3.588444e+04
min	370.000000	6.380000e+02
25%	1460.000000	5.000750e+03
50%	1980.000000	7.683000e+03
75%	2620.000000	1.100125e+04
max	13540.000000	1.074218e+06

In [38]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4600 entries, 0 to 4599
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   sqft_living  4600 non-null   int64
 1   sqft_lot    4600 non-null   int64
dtypes: int64(2)
memory usage: 72.0 KB
```

In [39]:

df.fillna(method='ffill',inplace=True)

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\4116506308.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.fillna(method='ffill',inplace=True)

In [40]:

```
x=np.array(df['sqft_living']).reshape(-1,1)
y=np.array(df['sqft_lot']).reshape(-1,1)
```

In [41]:

```
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```


0.054163304875602614

C:\Users\sruth\AppData\Local\Temp\ipykernel_18948\693062840.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
df.dropna(inplace=True)


In [42]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_pred,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [43]:

```
df500=df[:][:500]
sns.lmplot(x="sqft_living",y="sqft_lot",data=df500,order=1,ci=None)
x=np.array(df500['sqft_living']).reshape(-1,1)
y=np.array(df500['sqft_lot']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```




In [44]:

```
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
```

Regression: 0.07085964549127033

In [45]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In []:

In []:

In []: